

CPSC 427: Object-Oriented Programming

Michael J. Fischer

Lecture 24
December 3, 2018

Remarks about PS6

Circular dependencies

There is a circular dependency between classes `Block` and `SPtr`.

The following steps will break the circularity and avoid compiler errors.

1. Have `Block.hpp` `#include SPtr.hpp` but not the other way around, i.e., `SPtr.hpp` should **not** `#include Block.hpp`.
2. To fix the resulting error about `Block` being unknown when reading `SPtr.hpp`, one should add a forward declaration for `Block` at the top of `SPtr.hpp`. This looks like `class Block;`.
3. To compile `SPtr.cpp`, `#include Block.hpp` in the **CPP** file, `SPtr.cpp`.

Double-delete problem with SPtrs

All smart pointer pointing at the same **target** and **count** objects form a **smart pointer cluster**.

To form a cluster, a smart pointer should be constructed at the same time as a **Block** is constructed. Additional smart pointers to that same block should be created by copying a smart pointer that is already in the cluster for that block.

If a second cluster is created that points to the same block, the block will be deleted once for each cluster when it goes away, resulting in a double delete.

Using Blockchain objects

`Blockchain` is a wrapper around an `SPtr` object.

Blockchains can be freely assigned and copied using the default assignment and copy constructors.

The default assignment and copy constructor for `Blockchain` operates by assigning or copying corresponding data members.

For the `SPtr` data member, the assignment operator or copy constructor defined in `SPtr` will get used.

This results in the source and destination blockchains having embedded smart pointers that belong to the same cluster.

Nothing special has to be done in class `Blockchain` to achieve this behavior.

Clocks and Time Measurement

How to measure run time of a program

- ▶ There is no standard procedure in C++ for accurately measuring time.
- ▶ Time measurement depends on the software clocks provided by your computer and operating system.
- ▶ Clocks advance in discrete clicks called **jiffies**. A jiffy on the Zoo linux machines is one millisecond (0.001 seconds) long.
- ▶ Even if the clock is 100% accurate, the measured time can be off by as much as one jiffy.
- ▶ Hence, times shorter than tens of milliseconds cannot be directly measured with much accuracy using the standard software clock.

High resolution clocks

- ▶ Linux also provides high resolution clocks based on CPU timers.
- ▶ High resolution clocks are useful to the operating system for task scheduling and timeouts.
- ▶ They are also available to the user for higher-precision time measurements.
- ▶ Be aware that reading the clock involves a kernel call that takes a certain amount of time. This itself may limit the accuracy of timing measurements, even when the clock resolution is sufficiently high for the desired accuracy.
- ▶ See [man 7 time](#) for more information about linux clocks.

Measuring time in real systems

- ▶ Measuring code efficiency in real systems is challenging. Many factors can influence the results that are hard to control.
 - ▶ Other process running on the same machine.
 - ▶ Time spent in the OS moving data on and off disks.
 - ▶ Memory caching behavior.
- ▶ Lacking a controlled laboratory environment, one can still take steps to improve accuracy of tests:
 - ▶ Do some tests to determine what factors seem to have a sizable effect on the run time, e.g., the first run of a program is likely to be slower than subsequent runs because of caching.
 - ▶ Run the same test several times to get a feeling for the variance of results.
 - ▶ Make sure the optimizer isn't optimizing away code that you think is being executed.

Demo: Stopwatch

Realtime measurements

`StopWatch` is a class I wrote for measuring realtime performance of code.

It emulates a stopwatch with 3 buttons: `reset`, `start`, and `stop`.

At any time, the watch displays the cumulative time that the stopwatch has been running.

HirezTime class

`HirezTime` is a **wrapper class** for the system-specific functions to read the clock.

It hides the details of the underlying time representation and provides a simple interface for reading, computing, and printing times and time intervals.

`HirezTime` objects are intended to be copied rather than pointed at, and they try to behave like other numeric types.

Versions of HirezTime

There are two versions:

24-StopWatch (Linux/Unix/MacOSX) Function `gettimeofday()` returns the clock in a `struct timeval`, which consists of two `long ints` representing seconds and `microseconds`. The resolution of the clock is system-dependent, typically 1 millisecond. (See demo [24-StopWatch](#).)

24-StopWatch-hirez (Linux only) Function `clock_gettime()` returns the clock in a `struct timespec`, which consists of two `long ints` representing seconds and `nanoseconds`. The resolution of the clock is system-dependent and can be obtained with the `clock_getres()` function. (See demo [24-StopWatch-hirez](#).)

HirezTime structure

- ▶ In C++, `struct T` and `class T` are very similar. In both cases, `T` becomes a new type name.
- ▶ `struct` members are public by default.
`class` members are private by default.
- ▶ `HirezTime` is derived from `struct timeval` or `struct timespec`, depending on the version.
- ▶ It uses `protected` derivation to hide the underlying representation.
- ▶ It presents two interfaces to the world:
 1. The normal public interface treats `HirezTime` as an opaque object.
 2. A class derived from it can access the fields of the underlying `timespec/timeval`.

Printing a `HirezTime` number

Something seemingly simple like printing `HirezTime` values is not so simple. Naively, one might write:

```
cout << t.tv_sec << "." << t.tv_usec;
```

where `tv_sec` and `tv_usec` are the seconds and microseconds fields of a `timeval` structure.

If `t` represents 2 seconds and 27 microseconds, then what would print is `2.27`, not the correct `2.000027`.

The class contains a `print` function that fixes this problem.

StopWatch class

StopWatch contains five member variables to remember

- ▶ Whether the watch is running or not.
- ▶ The cumulative run time to point when last stopped.
- ▶ The most recent start and stop times.

All functions are `inline` to minimize inaccuracies of measurement due to the overhead within the stopwatch code itself.

Casting a Stopwatch to a HirezTime

An operator extension defines a cast for reading the cumulative time from a stop watch:

```
operator HirezTime() const { return cumSpan; }
```

Thus, if `sw` is a `StopWatch`,

```
cout << sw;
```

will print `sw.cumSpan` using `sw.print()`.

Why it works

This works because `operator<<()` is not defined for righthand operands of type `StopWatch` but it is defined for `HirezTime`.

The compiler then **coerces** `sw` to something that is acceptable to the `<<` operator.

Because `operator HiriezTime()` is defined for class `StopWatch`, the compiler will invoke it to obtain a `HirezTime` object, for which `<<` is defined.

Note that a similar coercion happens when one writes

```
if(!in) {...}
```

to test if an `istream` object `in` is open for reading. Here, the `istream` object is coerced to a `bool` because `operator bool()` is defined inside the `streams` package.